

## VALUE ITERATION

Exp No: 1

Date: 18.09.2024

### AIM:

To implement the dynamic Value Iteration programming algorithm to find the optimal policy and value function for a Markov Decision Process (MDP) in python.

### ALGORITHM:

#### Input:

- A set of states  $S$ .
- A set of actions  $A$ .
- Transition probabilities  $P(s'|s,a)$  — the probability of reaching state  $s'$  from state  $s$  by taking action  $a$ .
- Reward function  $R(s,a)$  — the reward obtained by taking action  $a$  in state  $s$ .
- Discount factor  $\gamma$  (gamma) — used to weigh the importance of future rewards.
- Convergence threshold  $\theta$  (theta) — used to determine when to stop iterations.

#### Output:

- Optimal value function  $V$  — the maximum expected reward for each state.
- Optimal policy  $\pi$  — the best action to take in each state.

#### Steps:

##### 1. Initialization:

- Initialize the value function  $V(s)$  for all states to 0.
- Set a convergence threshold  $\theta$ .

##### 2. Iterative Update:

- Repeat until the value function  $V$  converges:
  - Set  $\Delta = 0$ , which tracks the maximum change in the value function across all states.
  - For each state  $s$ :
    - Store the old value  $v = V(s)$ .

- For each action  $a$ , calculate the action-value  $Q(s,a)$  by summing over all possible next states  $s'$ :  $Q(s,a) = \sum_{s'} P(s'|s,a) \times [R(s,a) + \gamma \times V(s')]$   
 $Q(s,a) = \sum_{s'} P(s'|s,a) \times [R(s,a) + \gamma \times V(s')]$
- Update the value function  $V(s)$  with the maximum action-value:  
 $V(s) = \max_a Q(s,a)$
- Compute the difference between the new value and the old value:  
 $\Delta = \max_s (|V(s) - V_{old}(s)|)$
- If the maximum change  $\Delta$  is less than the threshold  $\theta$ , exit the loop, as the values have converged.

### 3. Policy Extraction:

- For each state  $s$ , compute the action-value  $Q(s,a)$  for each action  $a$  and derive the optimal action:  $\pi(s) = \arg\max_a Q(s,a)$
- The optimal policy  $\pi(s)$  assigns the best action for each state.

### 4. Return Results:

- Return the final value function  $V$  and the optimal policy  $\pi$ .

## SOURCE CODE:

```
import numpy as np
```

```
def value_iteration(states, actions, transition_prob, rewards, gamma=0.9, theta=1e-6):
    """ Performs Value Iteration for a Markov Decision Process (MDP)
    :param states: List of states
    :param actions: List of actions
    :param transition_prob: Transition probability (P(s'|s,a)) 3D array of shape (states, actions, states)
    :param rewards: Rewards as a 2D array of shape (states, actions)
    :param gamma: Discount factor
    :param theta: Convergence threshold
    :return: Optimal values (V) and optimal policy (pi) """

    # Initialize value function for each state
    V = np.zeros(len(states))
```

```
while True:
    delta = 0
    # Update each state's value
    for s in range(len(states)):
        v = V[s]
        Q_sa = np.zeros(len(actions))

        # Compute Q-value for each action
        for a in range(len(actions)):
            Q_sa[a] = sum([transition_prob[s, a, s_prime] * (rewards[s, a] + gamma * V[s_prime])
                           for s_prime in range(len(states))])

        # Update the value of the state with the best Q-value
        V[s] = max(Q_sa)
        delta = max(delta, abs(v - V[s]))

    # If the values converge, stop the iteration
    if delta < theta:
        break

# Derive the optimal policy
policy = np.zeros(len(states), dtype=int)
for s in range(len(states)):
    Q_sa = np.zeros(len(actions))
    for a in range(len(actions)):
        Q_sa[a] = sum([transition_prob[s, a, s_prime] * (rewards[s, a] + gamma * V[s_prime])
                       for s_prime in range(len(states))])
    policy[s] = np.argmax(Q_sa)
return V, policy
```

```
# Example usage:
states = [0, 1, 2]
actions = [0, 1]

# Transition probabilities (P(s'|s,a)): for each (state, action, next_state)
transition_prob = np.array([
    [[0.8, 0.2, 0], [0.1, 0.9, 0]], # Transitions from state 0
    [[0.5, 0.5, 0], [0.2, 0.8, 0]], # Transitions from state 1
    [[0.0, 0.0, 1], [0.0, 0.0, 1]] # Transitions from state 2])

# Rewards for each (state, action)
rewards = np.array([
    [1, -1], # Rewards for state 0
    [-1, 1], # Rewards for state 1
    [0, 0]  # Rewards for state 2])

# Perform value iteration
V, policy = value_iteration(states, actions, transition_prob, rewards)
print("Optimal Values: ", V)
print("Optimal Policy: ", policy)
```

### OUTPUT:

```
Optimal Values:  [9.99999245  9.99999288  0.          ]
Optimal Policy:  [0 1 0]
```

### RESULT:

Thus we have successfully implemented the dynamic Value Iteration programming algorithm to find the optimal policy and value function for a Markov Decision Process (MDP) in python.

## VALUE ITERATION FOR BRIDGE GRID ENVIRONMENT

Exp No: 2a

Date: 18.09.2024

### AIM:

To implement the solution for One Grid Environment (OGE) using value iteration in python.

### ALGORITHM:

#### Inputs:

- **Grid Size:** (rows,columns)
- **Bridge Location:** Starting point of the bridge (row,col)
- **Bridge Length:** Length of the bridge
- **Goal State:** Target state (row,col)
- **Fall Penalty:** Reward for falling off the bridge (negative value)
- **Goal Reward:** Reward for reaching the goal (positive value)
- **Discount Factor:**  $\gamma$  ( $0 < \gamma < 1$ )
- **Convergence Threshold:**  $\theta$  (small positive value)

#### Outputs:

- **Optimal Value Function:**  $V$  for each state
- **Optimal Policy:**  $\pi$  for each state

#### Steps:

##### 1. Initialize Environment:

- Create transition probability matrix  $P$  of shape (nstates,nactions,nstates).
- Create reward matrix  $R$  of shape (nstates,nactions).
- Set all state values  $V(s)=0$ .

##### 2. Build Environment:

- For each state (row,col):
  - For each action a:
    - If the current state is the goal state, set reward and transition.
    - Determine the next state based on the action:
      - If moving off the bridge, apply fall penalty.
      - Update the transition probabilities and rewards accordingly.

### 3. Value Iteration Loop:

- Repeat until convergence:
  - Set  $\delta=0$  (maximum change in value).
  - For each state  $s$ :
    - Store the current value:  $v=V(s)$ .
    - Calculate Q-values for all actions  $a$ :  $Q(s,a)=s'\sum P(s,a,s')\cdot(R(s,a)+\gamma\cdot V(s'))$
    - Update the value of state  $s$ :  $V(s)=\text{amax}Q(s,a)$
    - Update  $\delta$ :  $\delta=\max(\delta,|v-V(s)|)$
  - If  $\delta<\theta$ , stop the iteration.

### 4. Extract Optimal Policy:

- For each state  $s$ :
  - Calculate Q-values again for all actions.
  - Set policy:  $\pi(s)=\text{argamax}Q(s,a)$

### 5. Return:

- The optimal value function  $V$  and the optimal policy  $\pi$ .

## SOURCE CODE:

```
import numpy as np
```

```
class BridgeGridEnv:
```

```
    def __init__(self, grid_size=(5, 5), bridge_location=(2, 0), bridge_length=3, goal_state=(4, 4),  
fall_penalty=-10, goal_reward=10):
```

```
        self.grid_size = grid_size
```

```
        self.bridge_location = bridge_location
```

```
        self.bridge_length = bridge_length
```

```
        self.goal_state = goal_state
```

```
        self.fall_penalty = fall_penalty
```

```
        self.goal_reward = goal_reward
```

```
        self.actions = ['up', 'down', 'left', 'right']
```

```
self.n_states = grid_size[0] * grid_size[1]
```

```
self.n_actions = len(self.actions)
```

```
self.transitions, self.rewards = self.build_environment()
```

```
def build_environment(self):
```

```
    """Build transition and reward matrices for the environment."""
```

```
    transition_prob = np.zeros((self.n_states, self.n_actions, self.n_states))
```

```
    rewards = np.full((self.n_states, self.n_actions), -1) # Default reward for all actions is -1 (to  
encourage faster completion)
```

```
def state_to_index(row, col):
```

```
    return row * self.grid_size[1] + col
```

```
def is_off_bridge(row, col):
```

```
    return (row != self.bridge_location[0] or col >= self.bridge_location[1] + self.bridge_length)  
and row < self.grid_size[0] - 1
```

```
for row in range(self.grid_size[0]):
```

```
    for col in range(self.grid_size[1]):
```

```
        state = state_to_index(row, col)
```

```
        for action_idx, action in enumerate(self.actions):
```

```
            if (row, col) == self.goal_state:
```

```
                rewards[state, action_idx] = self.goal_reward
```

```
                transition_prob[state, action_idx, state] = 1.0
```

```
                continue
```

```
            next_row, next_col = row, col
```

```
            if action == 'up' and row > 0:
```

```
                next_row = row - 1
```

```
            elif action == 'down' and row < self.grid_size[0] - 1:
```

```
        next_row = row + 1
    elif action == 'left' and col > 0:
        next_col = col - 1
    elif action == 'right' and col < self.grid_size[1] - 1:
        next_col = col + 1

    next_state = state_to_index(next_row, next_col)

    if (next_row, next_col) == self.goal_state:
        rewards[state, action_idx] = self.goal_reward
    elif is_off_bridge(next_row, next_col):
        rewards[state, action_idx] = self.fall_penalty

    transition_prob[state, action_idx, next_state] = 1.0

    return transition_prob, rewards

def state_index_to_coordinates(self, state_index):
    """Convert state index back to row, col coordinates."""
    row = state_index // self.grid_size[1]
    col = state_index % self.grid_size[1]
    return row, col

def value_iteration_bridge(env, gamma=0.9, theta=1e-6):
    """ Perform value iteration to solve the Bridge Grid environment.
    """
    V = np.zeros(env.n_states)

    while True:
        delta = 0
```



```
for s in range(env.n_states):
    v = V[s]
    Q_sa = np.zeros(env.n_actions)

    # Calculate Q-values for all actions in state s
    for a in range(env.n_actions):
        Q_sa[a] = sum([env.transitions[s, a, s_prime] * (env.rewards[s, a] + gamma * V[s_prime])
                        for s_prime in range(env.n_states)])

    # Update state value with the maximum Q-value
    V[s] = max(Q_sa)
    delta = max(delta, abs(v - V[s]))

if delta < theta:
    break

# Derive policy from optimal state values
policy = np.zeros(env.n_states, dtype=int)
for s in range(env.n_states):
    Q_sa = np.zeros(env.n_actions)
    for a in range(env.n_actions):
        Q_sa[a] = sum([env.transitions[s, a, s_prime] * (env.rewards[s, a] + gamma * V[s_prime])
                        for s_prime in range(env.n_states)])
    policy[s] = np.argmax(Q_sa)

return V, policy

# Create the bridge grid environment
env = BridgeGridEnv()
```

```
# Run value iteration
V, policy = value_iteration_bridge(env)

# Display results
print("Optimal Values (V):")
print(V.reshape(env.grid_size))

print("\nOptimal Policy (state indices correspond to actions):")
print(policy.reshape(env.grid_size))
```

### OUTPUT:

```
Optimal Values (V):
[[27.70775002 32.89750102 38.66389102 37.78099102 45.79999102]
 [41.89750102 47.66389102 54.07099102 53.08999102 61.99999102]
 [47.66389102 54.07099102 61.18999102 70.09999102 79.99999102]
 [62.17099102 70.18999102 79.09999102 88.99999102 99.99999102]
 [70.18999102 79.09999102 88.99999102 99.99999102 99.99999102]]

Optimal Policy (state indices correspond to actions):
[[1 1 1 1 1]
 [1 1 1 1 1]
 [3 3 1 1 1]
 [1 1 1 1 1]
 [3 3 3 3 0]]
```

### RESULT:

Thus, the solution for One Grid Environment (OGE) using value iteration is implemented in Python.

## VALUE ITERATION FOR TWO GRID ENVIRONMENT

Exp No: 2b

Date: 18.09.2024

### AIM:

To implement the solution for Two Grid Environment (TGE) using value iteration in Python.

### ALGORITHM:

#### 1. Initialize Environment (BridgeGrid):

- Define a 5x5 grid.
- Set the goal position at (0, 4) (top-right corner).
- Set the start position at (4, 0) (bottom-left corner).
- Define a bridge of safe cells at positions (2, 1), (2, 2), (2, 3) on the middle row.
- Define possible actions as ['up', 'down', 'left', 'right'].
- Implement the reset, step, and render functions for the environment.

#### 2. Environment Dynamics (step function):

- **Actions:** Move the agent based on the chosen action (up, down, left, right) with boundary checks.
- **Rewards:**
  - Reward of +1 for reaching the goal.
  - Reward of 0 for being on the bridge.
  - Penalty of -1 for falling off the bridge into the water (outside the bridge in the middle row).
  - Small penalty of -0.1 for regular moves to encourage faster learning.
- Update the agent's state after the action.

#### 3. Initialize Q-Learning Agent (QLearningAgent):

- Set learning rate  $\alpha = 0.1$ , discount factor  $\gamma = 0.9$ , and exploration rate  $\epsilon = 0.1$ .
- Use a Q-table (implemented as a dictionary) to store Q-values for state-action pairs.

#### 4. Q-Learning Algorithm:

- **Q-value update formula:**  $Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma \max_a Q(s',a) - Q(s,a)]$   
 $Q(s,a) \leftarrow Q(s,a) + \alpha [r + \gamma \max_a Q(s',a) - Q(s,a)]$
- For each state-action pair, update the Q-value using the above formula, where:
  - s is the current state.
  - a is the current action.
  - r is the reward received.
  - s' is the next state after taking action a.
  - $\alpha$  is the learning rate, and  $\gamma$  is the discount factor.

### 5. Choose Action (choose\_action function):

- With probability  $\epsilon$ , choose a random action (exploration).
- With probability  $1 - \epsilon$ , choose the action that has the highest Q-value for the current state (exploitation).
- In case of ties (multiple actions with the same Q-value), randomly choose one of the tied actions.

### 6. Training Loop (train function):

- For each episode:
  - . **Reset** the environment to the initial state.
  - . Loop until the episode terminates (done = True):
    - Choose an action using the  $\epsilon$ -greedy strategy.
    - Take the action and observe the next state and reward.
    - Update the Q-value for the state-action pair.
    - Set the next state as the current state.
  - . If the episode is a multiple of 100, print progress.
- Repeat for the specified number of episodes (e.g., 1000).

### 7. Testing (after training):

- Reset the environment to the starting state.
- Choose actions based on the learned Q-values (exploitation).
- Render the grid after each step to visualize the agent's movement.
- Output the action taken and the corresponding reward after each step.

## SOURCE CODE:

```
import numpy as np
import random

# Environment Definition: Bridge Grid
class BridgeGrid:
    def __init__(self):
        self.grid = np.zeros((5, 5))
        self.goal = (0, 4) # Goal is at the top right corner
        self.start = (4, 0) # Start is at the bottom left corner
        self.bridge = [(2, 1), (2, 2), (2, 3)] # Narrow bridge in the middle row
        self.state = self.start
        self.actions = ['up', 'down', 'left', 'right']

    def reset(self):
        self.state = self.start
        return self.state

    def step(self, action):
        row, col = self.state

        if action == 'up':
            row = max(0, row - 1)
        elif action == 'down':
            row = min(4, row + 1)
        elif action == 'left':
            col = max(0, col - 1)
        elif action == 'right':
            col = min(4, col + 1)
```

```
next_state = (row, col)

# Define reward
if next_state == self.goal:
    reward = 1 # Reward for reaching the goal
    done = True
elif next_state in self.bridge:
    reward = 0 # Reward for being on the bridge
    done = False
elif row == 2 and col not in [1, 2, 3]:
    reward = -1 # Falling off the bridge (penalty)
    done = True
else:
    reward = -0.1 # Small penalty for each move to encourage faster learning
    done = False

self.state = next_state
return next_state, reward, done
```

```
def render(self):
    grid_copy = np.copy(self.grid)
    grid_copy[self.goal] = 1
    grid_copy[self.start] = 0.5
    for bridge_part in self.bridge:
        grid_copy[bridge_part] = 0.3
    print(grid_copy)
```

```
# Q-Learning Agent
class QLearningAgent:
```

```
def __init__(self, env, alpha=0.1, gamma=0.9, epsilon=0.1):
    self.env = env
    self.alpha = alpha # Learning rate
    self.gamma = gamma # Discount factor
    self.epsilon = epsilon # Exploration rate
    self.q_table = { } # Q-values dictionary

def get_q_value(self, state, action):
    return self.q_table.get((state, action), 0.0)

def update_q_value(self, state, action, reward, next_state):
    max_next_q_value = max([self.get_q_value(next_state, a) for a in self.env.actions])
    current_q_value = self.get_q_value(state, action)
    new_q_value = current_q_value + self.alpha * (reward + self.gamma * max_next_q_value -
current_q_value)
    self.q_table[(state, action)] = new_q_value

def choose_action(self, state):
    if random.uniform(0, 1) < self.epsilon:
        return random.choice(self.env.actions) # Explore
    else:
        # Exploit: choose action with the highest Q-value
        q_values = [self.get_q_value(state, action) for action in self.env.actions]
        max_q = max(q_values)
        max_actions = [action for action, q_value in zip(self.env.actions, q_values) if q_value ==
max_q]
        return random.choice(max_actions)

def train(self, episodes=1000):
    for episode in range(episodes):
        state = self.env.reset()
```

```
done = False
while not done:
    action = self.choose_action(state)
    next_state, reward, done = self.env.step(action)
    self.update_q_value(state, action, reward, next_state)
    state = next_state
if episode % 100 == 0:
    print(f"Episode {episode} complete")
```

```
# Create the environment
```

```
env = BridgeGrid()
```

```
# Create the Q-Learning agent
```

```
agent = QLearningAgent(env)
```

```
# Train the agent
```

```
agent.train(episodes=1000)
```

```
# Test the agent
```

```
state = env.reset()
```

```
env.render()
```

```
done = False
```

```
while not done:
```

```
    action = agent.choose_action(state)
```

```
    next_state, reward, done = env.step(action)
```

```
    state = next_state
```

```
    env.render()
```

```
    print(f"Action: {action}, Reward: {reward}\n")
```



**OUTPUT:**

Episode 0 complete  
Episode 100 complete  
Episode 200 complete  
Episode 300 complete  
Episode 400 complete  
Episode 500 complete  
Episode 600 complete  
Episode 700 complete  
Episode 800 complete  
Episode 900 complete

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

Action: up, Reward: -0.1

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

Action: right, Reward: -0

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

Action: up, Reward: 0

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

Action: right, Reward: 0

```
[[0. 0. 0. 0. 1. ]  
 [0. 0. 0. 0. 0. ]  
 [0. 0.3 0.3 0.3 0. ]  
 [0. 0. 0. 0. 0. ]  
 [0.5 0. 0. 0. 0. ]]
```

Action: right, Reward: 0

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: up, Reward: -0.1

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: left, Reward: -0.1

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: down, Reward: 0

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: right, Reward: 0

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: up, Reward: -0.1

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: up, Reward: -0.1

```
[[0. 0. 0. 0. 1. ]
 [0. 0. 0. 0. 0. ]
 [0. 0.3 0.3 0.3 0. ]
 [0. 0. 0. 0. 0. ]
 [0.5 0. 0. 0. 0. ]]
```

Action: right, Reward: 1

## RESULT:

Thus we implement the solution for Two Grid Environment (TGE) using value iteration in Python.

## VALUE ITERATION FOR DISCOUNT GRID ENVIRONMENT

Exp No: 3

Date: 18.09.2024

### AIM:

To implement the solution for Discount Grid Environment using value iteration in python.

### ALGORITHM

#### 1. Initialize the Discount Grid Environment (DiscountGridEnv):

- Create a grid environment with size (rows, cols) (default is 4x4).
- Set a goal state and assign a high reward for reaching the goal.
- Set penalty for each step (negative reward), and discount factor gamma for future rewards.
- Specify obstacle states where the agent cannot move.
- Define possible actions: up, down, left, right.

#### 2. Define Environment Dynamics (build\_environment):

- Build transition probabilities and reward structures for each state-action pair.
- For each state (row, col):
  - If it is the goal state, set the reward and transition to itself.
  - For other states, calculate the next state based on the action.
  - Ensure the agent doesn't move into obstacles or out of bounds.
- Return the transition probability matrix  $P(s, a, s')$  and reward function  $R(s, a)$ .

#### 3. Value Iteration Algorithm (value\_iteration\_discount\_grid):

- **Initialize** the value function  $V(s)$  for all states as 0.
- **Convergence criteria:** Use a small threshold theta to determine when to stop.

#### 4. Iteration Steps:

- For each state  $s$ , calculate the Q-value for all actions  $a$ :  
$$Q(s,a) = \sum_{s'} P(s,a,s') \times (R(s,a) + \gamma \times V(s'))$$
$$Q(s,a) = \sum_{s'} P(s,a,s') \times (R(s,a) + \gamma \times V(s'))$$
- Update the value  $V(s)$  of state  $s$  as the maximum Q-value across all actions:  
$$V(s) \leftarrow \max_a Q(s,a)$$

- Track the maximum difference between the old and new value functions (delta) to check convergence.
- If delta is smaller than the threshold theta, stop the iteration.

### 5. Derive Optimal Policy:

- After convergence of the value function, derive the optimal policy:
  - For each state s, select the action that maximizes the Q-value:  
$$\pi(s) = \arg\max_a Q(s, a)$$

### 6. Return the Optimal Value Function and Policy:

- Output the optimal value function  $V(s)$  and the optimal policy  $\pi(s)$ .

### SOURCE CODE:

```
import numpy as np
```

```
class DiscountGridEnv:
```

```
    def __init__(self, grid_size=(4, 4), goal_state=(3, 3), discount_factor=0.9, obstacle_states=[], goal_reward=10, step_penalty=-1):
```

```
        """
```

```
        Initializes the Discount Grid Environment.
```

```
        :param grid_size: The size of the grid (rows, cols)
```

```
        :param goal_state: The coordinates of the goal state
```

```
        :param discount_factor: Discount factor (gamma) for future rewards
```

```
        :param obstacle_states: List of obstacle coordinates (rows, cols) that the agent cannot pass through
```

```
        :param goal_reward: Reward for reaching the goal
```

```
        :param step_penalty: Penalty for taking a step (negative reward)
```

```
        """
```

```
        self.grid_size = grid_size
```

```
        self.goal_state = goal_state
```

```
        self.discount_factor = discount_factor
```

```
        self.obstacle_states = obstacle_states
```

```
        self.goal_reward = goal_reward
```

```
self.step_penalty = step_penalty
self.actions = ['up', 'down', 'left', 'right']
self.n_states = grid_size[0] * grid_size[1]
self.n_actions = len(self.actions)
self.transitions, self.rewards = self.build_environment()
```

```
def build_environment(self):
```

```
    """Build transition probabilities and rewards for each state-action pair.
    """
```

```
    transition_prob = np.zeros((self.n_states, self.n_actions, self.n_states))
```

```
    rewards = np.full((self.n_states, self.n_actions), self.step_penalty) # Default penalty for every
step
```

```
def state_to_index(row, col):
```

```
    return row * self.grid_size[1] + col
```

```
def is_valid_state(row, col):
```

```
    return (0 <= row < self.grid_size[0]) and (0 <= col < self.grid_size[1]) and (row, col) not in
self.obstacle_states
```

```
for row in range(self.grid_size[0]):
```

```
    for col in range(self.grid_size[1]):
```

```
        state = state_to_index(row, col)
```

```
        for action_idx, action in enumerate(self.actions):
```

```
            if (row, col) == self.goal_state:
```

```
                rewards[state, action_idx] = self.goal_reward
```

```
                transition_prob[state, action_idx, state] = 1.0
```

```
            continue
```

```
        next_row, next_col = row, col
```

```
if action == 'up' and is_valid_state(row - 1, col):
    next_row = row - 1
elif action == 'down' and is_valid_state(row + 1, col):
    next_row = row + 1
elif action == 'left' and is_valid_state(row, col - 1):
    next_col = col - 1
elif action == 'right' and is_valid_state(row, col + 1):
    next_col = col + 1
```

```
next_state = state_to_index(next_row, next_col)
```

```
transition_prob[state, action_idx, next_state] = 1.0
```

```
return transition_prob, rewards
```

```
def state_index_to_coordinates(self, state_index):
```

```
    """Convert state index back to (row, col) coordinates."""
```

```
    row = state_index // self.grid_size[1]
```

```
    col = state_index % self.grid_size[1]
```

```
    return row, col
```

```
def value_iteration_discount_grid(env, gamma=0.9, theta=1e-6):
```

```
    """ Perform Value Iteration for the Discount Grid environment.
```

```
    :param env: The DiscountGridEnv environment
```

```
    :param gamma: Discount factor
```

```
    :param theta: Convergence threshold for value iteration
```

```
    :return: Optimal value function (V) and optimal policy
```

```
    """
```

```
    V = np.zeros(env.n_states)
```

```
while True:
    delta = 0
    for s in range(env.n_states):
        v = V[s]
        Q_sa = np.zeros(env.n_actions)

        # Compute Q-values for all actions in state s
        for a in range(env.n_actions):
            Q_sa[a] = sum([env.transitions[s, a, s_prime] * (env.rewards[s, a] + gamma * V[s_prime])
                           for s_prime in range(env.n_states)])

        # Update the value of state s with the maximum Q-value
        V[s] = max(Q_sa)
        delta = max(delta, abs(v - V[s]))

    if delta < theta:
        break

# Derive policy from optimal state values
policy = np.zeros(env.n_states, dtype=int)
for s in range(env.n_states):
    Q_sa = np.zeros(env.n_actions)
    for a in range(env.n_actions):
        Q_sa[a] = sum([env.transitions[s, a, s_prime] * (env.rewards[s, a] + gamma * V[s_prime])
                       for s_prime in range(env.n_states)])
    policy[s] = np.argmax(Q_sa)
return V, policy

# Example usage
```

```
env = DiscountGridEnv(grid_size=(4, 4), goal_state=(3, 3), discount_factor=0.9,  
obstacle_states=[(1, 1)], goal_reward=10, step_penalty=-1)
```

```
# Run value iteration on the environment
```

```
V, policy = value_iteration_discount_grid(env, gamma=0.9)
```

```
# Display results
```

```
print("Optimal Value Function (V):")
```

```
print(V.reshape(env.grid_size))
```

```
print("\nOptimal Policy (actions corresponding to indices):")
```

```
policy_grid = np.array(env.actions)[policy].reshape(env.grid_size)
```

```
print(policy_grid)
```

### OUTPUT:

```
Optimal Value Function (V):  
[[48.45850102 54.95389102 62.17099102 70.18999102]  
 [54.95389102 62.17099102 70.18999102 79.09999102]  
 [62.17099102 70.18999102 79.09999102 88.99999102]  
 [70.18999102 79.09999102 88.99999102 99.99999102]]  
  
Optimal Policy (actions corresponding to indices):  
[['down' 'right' 'down' 'down']  
 ['down' 'down' 'down' 'down']  
 ['down' 'down' 'down' 'down']  
 ['right' 'right' 'right' 'up']]
```

### RESULT:

Thus, we have implemented the solution for Discount Grid Environment (DGE) using value iteration in Python.



## ACTION SELECTION METHOD

Exp No: 4

Date: 18.09.2024

### AIM:

To implement various action-selection strategies used in Reinforcement Learning (RL) to choose an action based on Q-values (which represent the estimated value of taking an action in a given state) for each state-action pair. To implement greedy, epsilon-greedy, softmax, and Upper Confidence Bound (UCB) strategies.

### ALGORITHM

#### 1. Greedy Action Selection

- **Input:** Q (2D array of state-action values), current state
- **Output:** Action with the highest value for the given state
- **Steps:**
  - . For the current state, look up the Q-values of all possible actions.
  - . Select the action that has the maximum Q-value.

#### 2. Epsilon-Greedy Action Selection

- **Input:** Q (2D array of state-action values), current state, epsilon (exploration probability)
- **Output:** Either a random action (with probability epsilon) or the greedy action (with probability  $1 - \epsilon$ )
- **Steps:**
  - . Generate a random number between 0 and 1.
  - . If the number is less than epsilon, select a random action (exploration).
  - . Otherwise, select the greedy action (action with the highest Q-value in the given state).

#### 3. Softmax Action Selection

- **Input:** Q (2D array of state-action values), current state, tau (temperature parameter controlling exploration)
- **Output:** Action selected probabilistically based on Q-values
- **Steps:**

- . For the current state, look up the Q-values of all possible actions.
- . Apply the softmax function to convert Q-values into a probability distribution over actions.
- . Select an action based on these probabilities (higher Q-values lead to higher probabilities, especially for lower  $\tau$ ).

#### 4. Upper Confidence Bound (UCB) Action Selection

- **Input:** Q (2D array of state-action values), current state, action counts, total counts, exploration parameter  $c$
- **Output:** Action selected based on both Q-value and exploration potential
- **Steps:**
  - . For each action in the given state:
    - If the action has never been taken, choose it immediately.
    - Otherwise, calculate the UCB value for the action using the formula:  

$$UCB(a) = Q(\text{state}, a) + c \times \sqrt{\frac{\log(\text{total counts})}{\text{action counts}(\text{state}, a)}}$$

$$UCB(a) = Q(\text{state}, a) + c \times \sqrt{\frac{\log(\text{total counts})}{\text{action counts}(\text{state}, a)}}$$
  - . Select the action with the highest UCB value.

#### SOURCE CODE:

```
import numpy as np
import random

def greedy_action_selection(Q, state):
    """
    Select the action with the highest value (greedy).

    :param Q: A 2D array where Q[s, a] is the estimated value of action 'a' in state 's'
    :param state: Current state
    :return: Selected action (greedy)
    """
    return np.argmax(Q[state])
```

```
def epsilon_greedy_action_selection(Q, state, epsilon=0.1):
    """
    Select an action using the epsilon-greedy strategy.
    :param Q: A 2D array where Q[s, a] is the estimated value of action 'a' in state 's'
    :param state: Current state
    :param epsilon: Probability of selecting a random action (exploration)
    :return: Selected action
    """
    if random.uniform(0, 1) < epsilon:
        # Exploration: Choose a random action
        return random.choice(range(Q.shape[1]))
    else:
        # Exploitation: Choose the action with the highest Q-value
        return np.argmax(Q[state])

def softmax_action_selection(Q, state, tau=1.0):
    """ Select an action using the softmax strategy.
    :param Q: A 2D array where Q[s, a] is the estimated value of action 'a' in state 's'
    :param state: Current state
    :param tau: Temperature parameter controlling exploration; high tau means more exploration
    :return: Selected action
    """
    q_values = Q[state]
    # Apply the softmax transformation to the Q-values
    exp_q = np.exp(q_values / tau)
    action_probabilities = exp_q / np.sum(exp_q)

    # Choose an action based on the computed probabilities
    return np.random.choice(len(q_values), p=action_probabilities)
```

```
def ucb_action_selection(Q, state, action_counts, total_counts, c=1.0):
    """ Select an action using the Upper Confidence Bound (UCB) strategy.
    :param Q: A 2D array where Q[s, a] is the estimated value of action 'a' in state 's'
    :param state: Current state
    :param action_counts: A 2D array tracking how many times each action has been taken in each
    state
    :param total_counts: Total number of actions taken so far
    :param c: Exploration parameter (higher means more exploration)
    :return: Selected action
    """
    ucb_values = np.zeros(Q.shape[1])

    for a in range(Q.shape[1]):
        if action_counts[state, a] == 0:
            return a # If action has never been taken, choose it
        else:
            ucb_values[a] = Q[state, a] + c * np.sqrt(np.log(total_counts) / (action_counts[state, a] + 1e-
5))

    return np.argmax(ucb_values)

# Example Q-values for a simple environment with 3 states and 4 actions per state
Q = np.array([
    [1.0, 0.5, 0.2, 0.8], # Q-values for state 0
    [0.1, 2.0, 0.3, 0.4], # Q-values for state 1
    [0.5, 0.4, 3.0, 1.0] # Q-values for state 2])

# Initialize action counts for UCB
action_counts = np.zeros(Q.shape) # Keeps track of how many times each action has been chosen
total_counts = 1 # Total action selections (initially 1 to avoid division by zero)
```

```
# Current state
state = 0

# Greedy action selection
action_greedy = greedy_action_selection(Q, state)
print(f"Greedy action selected: {action_greedy}")

# Epsilon-greedy action selection
action_epsilon_greedy = epsilon_greedy_action_selection(Q, state, epsilon=0.1)
print(f"Epsilon-greedy action selected: {action_epsilon_greedy}")

# Softmax action selection
action_softmax = softmax_action_selection(Q, state, tau=1.0)
print(f"Softmax action selected: {action_softmax}")

# UCB action selection
action_ucb = ucb_action_selection(Q, state, action_counts, total_counts, c=2.0)
print(f"UCB action selected: {action_ucb}")
```

### **OUTPUT:**

```
Greedy action selected: 0
Epsilon-greedy action selected: 0
Softmax action selected: 3
UCB action selected: 0
```

### **RESULT:**

Thus, we have successfully implemented various action-selection strategies used in Reinforcement Learning (RL) to choose an action based on Q-values (which represent the estimated value of taking an action in a given state) for each state-action pair. We have successfully implemented greedy, epsilon-greedy, softmax, and Upper Confidence Bound (UCB) strategies.

## Q-LEARNING ALGORITHM

Exp No: 5

Date: 18.09.2024

### AIM:

To implement a Q-learning algorithm in a grid world environment where an agent learns to navigate from the start to a goal while avoiding traps, using an epsilon-greedy strategy to balance exploration and exploitation which trains the agent by updating Q-values based on rewards and tests the learned policy to evaluate performance.

### ALGORITHM:

#### 1. Initialize the Environment (GridworldEnv Class):

- **Gridworld Environment Setup:**
  - Grid size: Defined by `grid_size` (default 4x4).
  - Goal state: Reaching this state gives a positive reward (`goal_reward`).
  - Trap states: Falling into these gives a penalty (`trap_penalty`).
  - Step penalty: Every move has a negative reward (`step_penalty`).
- **Agent Actions:** The agent can move up, down, left, or right.
- **Initial State:** The agent starts at the top-left corner (0, 0).
- **Step Function:** Moves the agent based on the action and checks if the agent has reached the goal, a trap, or just a normal position. Rewards and done flag are updated accordingly.

#### 2. Q-Learning Algorithm:

- **Initialize Q-Table:**
  - Q-table is a 3D array of size `grid_size[0] x grid_size[1] x n_actions`. Initially, all Q-values are set to zero.
- **For each episode:**
  - **Reset the Environment:** The agent starts at the initial position (0, 0).
  - **While the episode is not done:**
    - **Epsilon-Greedy Action Selection:**
      - With probability epsilon, choose a random action (exploration).

- With probability  $1 - \epsilon$ , choose the action with the highest Q-value for the current state (exploitation).
- **Take Action:** Apply the chosen action, observe the next state, reward, and whether the episode has ended (goal or trap reached).
- **Q-Value Update:**
  - Update the Q-value for the current state-action pair using the Q-learning update rule:  $Q[s,a] = (1 - \alpha)Q[s,a] + \alpha(r + \gamma \max_{a'}(Q[s',a']))$
  - Where:
    - $Q[s,a]$  is the current Q-value for the state-action pair.
    - $\alpha$  is the learning rate.
    - $r$  is the observed reward.
    - $\gamma$  is the discount factor for future rewards.
    - $\max_{a'}(Q[s',a'])$  is the highest Q-value for the next state  $s'$ .
- **Move to Next State:** Update the current state to the next state and repeat.
- **End of Episode:** Once the episode ends, move to the next episode and repeat the process.

### 3. Test the Learned Policy:

- **Using the Trained Q-Table:**
  - Start from the initial state (0, 0).
  - Always choose the action with the highest Q-value (greedy).
  - Execute the actions until the agent reaches the goal or a trap.
  - Output the steps taken, actions performed, and the total reward accumulated during the test.

### SOURCE CODE:

```
import numpy as np
import random
```

```
class GridworldEnv:

    def __init__(self, grid_size=(4, 4), goal_state=(3, 3), trap_states=[], goal_reward=10,
trap_penalty=-10, step_penalty=-1):

        """Initializes the Gridworld environment.

        :param grid_size: The size of the grid (rows, cols)
        :param goal_state: The coordinates of the goal state
        :param trap_states: List of trap coordinates (rows, cols) that give negative rewards
        :param goal_reward: Reward for reaching the goal
        :param trap_penalty: Penalty for falling into traps
        :param step_penalty: Penalty for taking a step (negative reward)
        """

        self.grid_size = grid_size
        self.goal_state = goal_state
        self.trap_states = trap_states
        self.goal_reward = goal_reward
        self.trap_penalty = trap_penalty
        self.step_penalty = step_penalty
        self.actions = ['up', 'down', 'left', 'right']
        self.n_actions = len(self.actions)

    def reset(self):

        """Reset the environment to the initial state (top-left corner)."""

        self.agent_position = (0, 0)
        return self.agent_position

    def step(self, action):

        """Take a step in the environment according to the action."""

        row, col = self.agent_position

        if action == 0: # up
```



```
        next_position = (max(row - 1, 0), col)
    elif action == 1: # down
        next_position = (min(row + 1, self.grid_size[0] - 1), col)
    elif action == 2: # left
        next_position = (row, max(col - 1, 0))
    elif action == 3: # right
        next_position = (row, min(col + 1, self.grid_size[1] - 1))

    reward = self.step_penalty
    done = False

    if next_position == self.goal_state:
        reward = self.goal_reward
        done = True
    elif next_position in self.trap_states:
        reward = self.trap_penalty
        done = True

    self.agent_position = next_position
    return next_position, reward, done

def get_state(self):
    """Returns the current state (row, col)."""
    return self.agent_position

def action_space(self):
    """Returns the number of available actions."""
    return self.n_actions

def q_learning(env, episodes=1000, alpha=0.1, gamma=0.9, epsilon=0.1):
```

```
"""Q-Learning algorithm implementation.
:param env: The environment
:param episodes: Number of episodes to run
:param alpha: Learning rate
:param gamma: Discount factor
:param epsilon: Exploration rate
:return: Learned Q-values
"""

# Initialize the Q-table (states: grid_size, actions: up, down, left, right)
Q = np.zeros((env.grid_size[0], env.grid_size[1], env.n_actions))

for episode in range(episodes):
    state = env.reset()
    done = False

    while not done:
        row, col = state

        # Epsilon-greedy action selection
        if random.uniform(0, 1) < epsilon:
            action = random.choice(range(env.n_actions))
        else:
            action = np.argmax(Q[row, col])

        # Take action, observe reward and next state
        next_state, reward, done = env.step(action)
        next_row, next_col = next_state

        # Update Q-value using the Q-learning update rule
        best_next_action = np.argmax(Q[next_row, next_col])
```

```
td_target = reward + gamma * Q[next_row, next_col, best_next_action]
Q[row, col, action] = (1 - alpha) * Q[row, col, action] + alpha * td_target
```

```
state = next_state
```

```
return Q
```

```
def test_q_learning(Q, env):
```

```
    """Test the learned Q-values by running the agent through the environment.
```

```
    :param Q: The learned Q-values
```

```
    :param env: The environment
```

```
    """
```

```
    state = env.reset()
```

```
    done = False
```

```
    steps = 0
```

```
    total_reward = 0
```

```
    print("Testing the learned policy...")
```

```
    while not done:
```

```
        row, col = state
```

```
        action = np.argmax(Q[row, col]) # Exploit the learned Q-values (greedy)
```

```
        next_state, reward, done = env.step(action)
```

```
        total_reward += reward
```

```
        steps += 1
```

```
        print(f"Step {steps}: State {state}, Action {env.actions[action]}, Reward {reward}")
```

```
        state = next_state
```

```
print(f"Total Reward: {total_reward}, Steps Taken: {steps}")

# Create the Gridworld environment

env = GridworldEnv(grid_size=(4, 4), goal_state=(3, 3), trap_states=[(1, 1)], goal_reward=10,
trap_penalty=-10)

# Run Q-Learning

Q = q_learning(env, episodes=1000, alpha=0.1, gamma=0.9, epsilon=0.1)

# Test the learned Q-values

test_q_learning(Q, env)
```

### OUTPUT:

```
Testing the learned policy...
Step 1: State (0, 0), Action down, Reward -1
Step 2: State (1, 0), Action down, Reward -1
Step 3: State (2, 0), Action right, Reward -1
Step 4: State (2, 1), Action down, Reward -1
Step 5: State (3, 1), Action right, Reward -1
Step 6: State (3, 2), Action right, Reward 10
Total Reward: 5, Steps Taken: 6
```

### RESULT:

Thus, we have successfully implemented a Q-learning algorithm in a grid world environment where an agent learns to navigate from the start to a goal while avoiding traps, using an epsilon-greedy strategy to balance exploration and exploitation which trains the agent by updating Q-values based on rewards and tests the learned policy to evaluate performance.